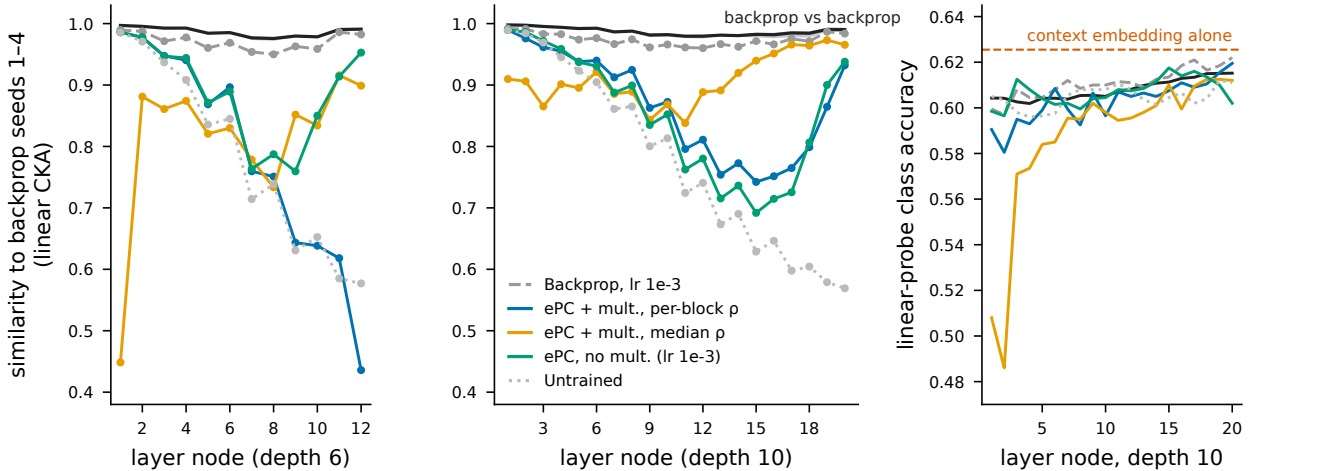
a
Depth 6: PC's middle and late layers differ from backprop
similarity to backprop seeds 1–4 (linear CKA)
layer node (depth 6)
b
Depth 10: PC differs in the middle, re-converges near the output
backprop vs backprop
Backprop, lr 1e-3
ePC + mult., per-block ρ
ePC + mult., median ρ
ePC, no mult. (lr 1e-3)
Untrained
layer node (depth 10)
c
Hidden states hold no more class information than the context
context embedding alone
linear-probe class accuracy
layer node, depth 10